Stan Lee Natividad

AI Infrastructure & Machine Learning Engineer

Mandaluyong City, Metro Manila 1441, PH · +63 991 153 2620 · leenatividad5@gmail.com · gryphlee.github.io

PROFILE

AI Engineer specializing in local, self-hosted AI systems — fine-tuning open-source LLMs (LoRA, QLoRA, SFT, knowledge distillation, GRPO/DPO) and deploying them behind production RAG pipelines and multi-agent orchestration. Strong focus on the parts of the lifecycle most teams skip: evaluation and observability (RAGAS, DeepEval; groundedness, hallucination, and latency tracking) and adversarial security testing (jailbreak and injection red-teaming). Backed by hands-on infrastructure experience across GPU compute, Ray distributed training, and vLLM serving.

EMPLOYMENT HISTORY

AI Infrastructure Engineer & Machine Learning Engineer | Kooya Inc *Jun 2026 — Present*

BGC, Taguig

- Manage the AI platform used by engineering teams to fine-tune and serve models — GPU compute allocation, distributed training with Ray, and high-performance inference via vLLM.
- Fine-tune open-source LLMs (LoRA, QLoRA, SFT, knowledge distillation, GRPO, DPO) end-to-end, from dataset preparation through evaluation and production-readiness review.
- Maintain and extend the company's production RAG pipeline and multi-agent orchestration, serving 4 internal teams and ~1,200 queries per day.
- Built observability across serving infrastructure, RAG pipelines, and local models using RAGAS and DeepEval — cutting measured hallucination rate from 12% to 4% and p95 latency from 2.1s to 850ms.
- Red-team in-house models before deployment (jailbreak and prompt/system-injection testing), surfacing and patching 15+ exploitable behaviors pre-launch.
- Built a reproducible benchmarking environment comparing open-source models on reasoning quality, latency, and cost — informing model selection for the company's internal knowledge assistant.
- Supporting an ongoing migration from third-party APIs to locally hosted models, projected to cut inference costs by ~60% (≈₱180K/month).

Intelligent Systems Engineer | Power4All (6-month project engagement) *Jan 2026 — Jun 2026*

Pasig

- Designed and delivered a fully licensed, company-only local AI model for the engineering and business development teams, replacing third-party AI services for document drafting, technical Q&A, and proposal support.
- Built a centralized company knowledge base with an end-to-end RAG pipeline, unifying company files and operational data into a single retrieval backbone.
- Automated the bid-discovery workflow for business development — auto-scraping opportunities from external sites into an internal platform, saving an estimated 10 hours per week of manual searching.
- Built a digital twin of the company's water-treatment plant, enabling the team to simulate chemical dosing in software before applying changes to the physical plant — reducing trial-and-error dosing adjustments by ~80%.
- Initiated broader Industry 4.0 work including AI-integrated SCADA monitoring and workflow automation across supply chain and engineering functions.

Enterprise AI Architect | Green Attic Roofing *Dec 2024 — Dec 2025*

Chicago (Remote)

- Converted manual, repetitive operations into automated AI-driven workflows across 3 departments, saving an estimated 25+ hours per week of staff time.
- Built internal AI tools on request from teams across the business — integrating LLMs, external APIs, and internal databases into unified workflows.
- Created reusable AI frameworks that cut build time for new automations from ~2 weeks to ~3 days.
- Served as the organization's AI adoption lead — evaluating tools and delivering solutions built for security and maintainability.

AI Engineer | Silicon Valley Associate *Jan 2024 — Nov 2024*

Remote (UAE-based consultancy)

- Built an AI-powered hiring platform with LLM-driven candidate matching, screening, and routing, processing 2,000+ applications.
- Designed and shipped an AI interviewer backed by a full RAG pipeline — retrieving role-specific context to drive dynamic questioning and evaluation, reducing time-to-screen by ~70%.
- Worked directly with the CEO to convert manual hiring workflows into scalable agentic systems.

EDUCATION

BS Electronics Engineering | Rizal Technological University *Mandaluyong*

Completed coursework through 3rd year

CORE SKILLS

AI & LLM: Fine-Tuning (LoRA, QLoRA, Adapters, SFT, GRPO, DPO) · Knowledge Distillation · Model Quantization · Function Calling & Tool Use · Local Model Deployment

RAG & Agents: End-to-End RAG Pipelines · Multi-Agent Systems · Agent Orchestration

Evaluation & Observability: RAGAS · DeepEval · Model & RAG Benchmarking · Groundedness, Faithfulness & Hallucination Tracking · Latency Monitoring

AI Security: Adversarial Testing & Red-Teaming · Jailbreak & Prompt-Injection Testing · Model Hardening

Infrastructure: GPU Clusters · Distributed Training (Ray) · vLLM Serving · PyTorch · HuggingFace · ONNX/TensorRT · Kubernetes (working knowledge)

Industrial & Automation: Industry 4.0 · Digital Twin Simulation · AI-SCADA · Workflow Automation (n8n, Zapier)